

ARUN BHAT

Senior Platform Engineer

+91 9964862970 | arun.b.bhat@gmail.com | [LinkedIn](#)

PROFILE

Certified Senior Platform Engineer (CCDAK, Data Streaming Engineer, CKA) with proven experience designing, building, and operating production data pipelines on Kafka-based streaming and CDC platforms that support enterprise analytics and operational systems. Skilled in implementing data ingestion, transformation, and storage solutions with a strong focus on reliability, performance, and data quality across standardized, scalable data platforms. Experienced collaborating with cross-functional teams and automating infrastructure with Python, Shell scripting, and Terraform to enable consistent, governed data consumption.

SKILLS

Data Pipelines • Apache Kafka • Confluent Platform • Debezium CDC • Kubernetes • Python • Shell Scripting • Terraform • Docker • DevOps • Grafana • OpenSearch • Ansible

EXPERIENCE

Senior Platform Engineer | iHerb

2025 – Present

- Operate enterprise-scale Confluent Kafka and Strimzi-based clusters on Kubernetes as part of the standardized data platform, ensuring scalability, reliability, and security.
- Provisioned dedicated Kafka clusters to support isolated, high-throughput workloads for critical data pipelines.
- Design and operate large-scale Debezium CDC data ingestion pipelines, including connector lifecycle, incident response, and reliability improvements.
- Improved CDC streaming throughput by ~60% through connector and task-level performance tuning.
- Reduced Confluent Kafka costs by streamlining topic and data retention governance removing unused topics, right-sizing partitions, and tuning retention settings.
- Implemented connector health and data-freshness monitoring and alerting, enabling proactive detection of data quality issues over reactive firefighting.
- Contributed to root-cause analysis of a production OOM and silent data-flow stall, helping diagnose a heap-to-container memory issue and an unbounded Debezium catch-up query, and fixing both via resizing and config tuning.
- Collaborated on designing a new alerting framework (event-drought, restart-rate, offset-stall, heap-pressure alerts) to catch silent Kafka Connect failures reporting a healthy state.
- Wrote a custom Java SMT (Single Message Transform) for the Trustpilot DynamoDB source connector, fixing a data-loss bug that silently dropped attribute types preserving data quality and completeness downstream
- Contributed to capacity planning and CKU sizing for a Confluent Cloud cluster consolidation, using a LogEndOffset-based method to identify active partitions on compacted topics.
- Participated in evaluating Confluent Cloud topology and networking options (Dedicated vs. Standard vs. Enterprise, AWS Transit Gateway) for a cost-effective, compliant architecture.
- Provisioned and managed Confluent Cloud clusters and other AWS services
- Work extensively on cloud-managed K8s clusters like EKS, using Helm, Kustomize, and Flux CD to manage GitOps-based deployments and configuration for Kafka components and Kubernetes workloads.
- Troubleshoot Kubernetes production issues — CrashLoopBackOff, OOMKilled pods, resource constraints, networking, DNS, and persistent storage.

Data Engineer / Lead SRE | Platformary Labs

2022 – 2025

Client: Jollibee Foods Corporation

- Contributed to the design of Fusion Data Hub (FDH), a centralized Kafka-based data hub for enterprise data integration.

- Provisioned and managed Kubernetes clusters to support platform and application workloads.
- Implemented CRM journeys using Python, orchestrated through Temporal workflows and deployed within a Docker Compose environment.
- Collaborated with 40+ engineering teams to plan integrations and provide recommendations for reliable, scalable data ingestion.
- Configured REST Proxy to support data movement and service communication within EKS clusters.

Client: Tata Digital

- Contributed to the strategic migration of streaming data pipelines from Azure Event Hubs to Confluent Cloud Kafka, ensuring a seamless transition with zero data loss.
- Collaborated on networking and security design for the migration, defining connectivity requirements, network policies, security groups, and access controls.
- Provisioned and managed Confluent Cloud Kafka infrastructure using Terraform, supporting repeatable and consistent environment deployments.
- Supported implementation of monitoring and observability dashboards for Kafka infrastructure and streaming pipelines.
- Achieved over 50% cost optimization through platform migration, infrastructure right-sizing, and efficient resource management.

Client: MCX (Multi Commodity Exchange)

- Led a team of six SREs responsible for 24/7 availability and operational reliability of regional Confluent Kafka clusters.
- Helped plan and coordinate rotational on-call shifts, ensuring continuous production coverage and effective incident ownership.
- Managed P1/P2 production incidents, coordinated troubleshooting and recovery efforts, and prepared detailed RCA reports to drive preventive improvements.
- Established a structured process for periodic platform health checks, proactively identifying Kafka, infrastructure, capacity, and service-level issues.
- Managed on-call alerts and incident response through PagerDuty, improving escalation, ownership, and response to critical production issues.
- Implemented an ELK stack for centralized log aggregation and automated deployment/configuration of Confluent Platform monitoring components.
- Participated in Confluent Platform upgrades on bare-metal infrastructure, using Ansible to automate controlled rolling upgrades in restricted, no-internet environments via jump hosts.
- Reduced processing time by 3–4 hours by transitioning from batch to real-time stream processing.

CERTIFICATIONS

- [Confluent Certified Developer for Apache Kafka \(CCDAK\)](#)
- [Certified Kubernetes Administrator \(CKA\)](#)
- [Confluent Certified Data Streaming Engineer](#)

EDUCATION

B.Tech in Instrumentation and Control Systems

Manipal Institute of Technology